

Recommendations on AWS FIS

AWS Fault Injection Simulator (FIS) is a powerful tool that allows you to conduct controlled experiments on your AWS infrastructure to uncover weaknesses, improve resiliency, and enhance the overall stability of your systems. Here are some recommendations for effectively using FIS in your chaos engineering practices:

1. **Start with Clear Objectives:** Define the goals and objectives of your chaos engineering experiments. Are you trying to identify single points of failure, test failover mechanisms, or assess the impact of resource failures? Having clear objectives will guide your experiments and help you derive meaningful insights.
2. **Identify Target Services:** Choose the AWS services and resources that you want to test. Focusing on critical components of your architecture, such as databases, load balancers, or compute instances, will provide valuable insights into your system's behavior under stress.
3. **Create Experiment Templates:** Experiment templates in FIS define the parameters of your experiments, such as the actions to be taken and the targets. Create templates that simulate scenarios like instance termination, network interruptions, or scaling events. This will ensure consistency and repeatability in your experiments.
4. **Start with Low Impact Experiments:** Begin with chaos experiments that have a relatively low impact on your production environment. This helps you gain confidence in the FIS tool and understand how your infrastructure responds to controlled failures before moving on to more impactful experiments.
5. **Gradually Increase Complexity:** As you become more comfortable with FIS, gradually increase the complexity and severity of your experiments. This could involve triggering multiple failures simultaneously or introducing failures in different regions to test cross-region redundancy.
6. **Monitor and Measure:** Set up comprehensive monitoring and observability tools to track the behavior of your systems during chaos experiments. Use CloudWatch metrics, CloudTrail logs, and any other relevant monitoring solutions to gain insights into how your infrastructure responds to failures.

7. **Document and Analyze Results:** Document the results of each experiment, including the impact on your application's performance, the behavior of AWS services, and any unexpected outcomes. Analyze the data to identify patterns, bottlenecks, and potential areas for improvement.
8. **Iterate and Improve:** Chaos engineering is an iterative process. Use the insights gained from each experiment to refine your infrastructure, architecture, and failover mechanisms. Implement changes based on lessons learned to enhance your system's resiliency.
9. **Collaborate and Communicate:** Involve your development, operations, and security teams in the chaos engineering process. Effective communication ensures that everyone understands the goals, potential risks, and benefits of chaos experiments, fostering a culture of resiliency.
10. **Practice Regularly:** Chaos engineering is not a one-time activity. Regularly conduct chaos experiments to ensure that your systems remain resilient as they evolve. New deployments, updates, and configuration changes can introduce vulnerabilities that might not be evident until tested.
11. **Consider Cost Implications:** Some chaos experiments may involve scaling resources or inducing failures that could have cost implications. Be mindful of potential cost increases and set up appropriate cost controls to prevent unexpected charges.

Remember that while chaos engineering is a powerful technique for enhancing system resilience, it should be performed thoughtfully and strategically to minimize disruption to your production environment. AWS FIS provides a controlled way to test failures, but it's important to balance the benefits of chaos experiments with the potential risks they might introduce.

How to use AWS Fault Injection Simulator (FIS):

1. Log in to your AWS Management Console using your credentials.
2. Navigate to the "Fault Injection Simulator" service from the list of available services in the AWS Management Console.
3. Click on "Experiment templates" in the left navigation pane.
4. Click the "Create template" button.
5. Give your template a name and description to describe its purpose.

6. In the "Actions" section, click "Add action."
7. Choose an action type.
8. Configure the action parameters, such as the target resources and duration of the disruption.
9. In the "Targets" section, click "Add target."
10. Choose the AWS resources you want to target for the experiment. This could be EC2 instances, Auto Scaling groups, or other services.
11. Configure the resource selection criteria based on tags, resource IDs, or other identifiers.
12. In the "Permissions" section, choose the IAM roles that FIS should assume to execute the experiment or create a new role.
13. Make sure the selected roles have the necessary permissions to modify the targeted resources.
14. Review your experiment template's configuration for accuracy.
15. Click the "Create" button to save the template.
16. Go back to the "Experiment templates" page.
17. Locate the template you created and click the "Run experiment" button.
18. Select the template version to run if you've made changes to it.
19. Configure any experiment-specific parameters, such as the duration and steady state period.
20. Once the experiment starts, monitor its progress in the FIS dashboard.
21. Use AWS CloudWatch metrics, CloudTrail logs, and any other monitoring tools to observe the impact of the experiment on your resources.
22. After the experiment is complete, review the results to understand how your resources and applications responded.
23. Analyze any unexpected behavior or performance issues that arose during the experiment.

Example Experiment Template to Test VPC Resiliency:

Let's consider an example of an experiment template to test the resiliency of a Virtual Private Cloud (VPC):

Template name: VPC-Resilience-Test

Description: Test the impact of network disruptions on a VPC.

Actions:

- Action Type: Network disruption
- Parameters: Disconnect TCP traffic
- Duration: 5 minutes
- Scope: vpc

Targets:

- Resource type: aws:ec2:subnet
- Resource selection: Based on subnets tags (e.g., "Environment:Production")

Permissions:

- Use an IAM role with permissions to modify VPC settings.

This experiment template simulates a network disruption by disconnecting TCP traffic in the targeted VPC for a duration of 5 minutes. The experiment can help you understand how your applications and services within the VPC respond to network failures and whether they gracefully recover after the disruption.

Remember that experiment templates and configurations will vary based on your specific use case and the objectives you want to achieve with AWS FIS. Always ensure that your experiments are well-planned and conducted in a controlled environment to minimize any impact on your production systems.